

Zeppelin对接FusionInsight HD

适用场景

Zeppelin 0.7.2 ↔ FusionInsight HD V100R002C60U20

安装Zeppelin

操作场景

安装Zeppelin0.7.2

前提条件

- 已完成FusionInsight HD客户端的安装。

操作步骤

- 将软件包zeppelin-0.7.2-bin-all.tgz上传至/opt目录下，解压生成zeppelin-0.7.2-bin-all目录。

```
tar -zxvf zeppelin-0.7.2-bin-all.tgz
```

- 启动和停止Zeppelin

```
bin/zeppelin-daemon.sh start
bin/zeppelin-daemon.sh stop
```

```
[root@localhost zeppelin-0.7.2-bin-all]# bin/zeppelin-daemon.sh start
Log dir doesn't exist, create /opt/zeppelin-0.7.2-bin-all/logs
Pid dir doesn't exist, create /opt/zeppelin-0.7.2-bin-all/run
Zeppelin start [ OK ]
[root@localhost zeppelin-0.7.2-bin-all]# bin/zeppelin-daemon.sh stop
Zeppelin stop [ OK ]
```

- 配置Zeppelin环境变量，在profile文件中加入如下变量

```
vi /etc/profile
export ZEPPELIN_HOME=/opt/zeppelin-0.7.2-bin-all
export PATH=$ZEPPELIN_HOME/bin:$PATH
```

- 编辑zeppelin-env.sh文件，位置/opt/zeppelin-0.7.2-bin-all/conf

```
cd /opt/zeppelin-0.7.2-bin-all/conf/
cp zeppelin-env.sh.template zeppelin-env.sh
vi zeppelin-env.sh
```

加入如下内容：

```
export JAVA_HOME=/opt/jdk1.7.0_51/
```

编辑zeppelin-site.xml文件，位置/opt/zeppelin-0.7.2-bin-all/conf/

```
cp zeppelin-site.xml.template zeppelin-site.xml
```

将zeppelin-site.xml中端口8080替换成18081（可自定义，也可以不改）

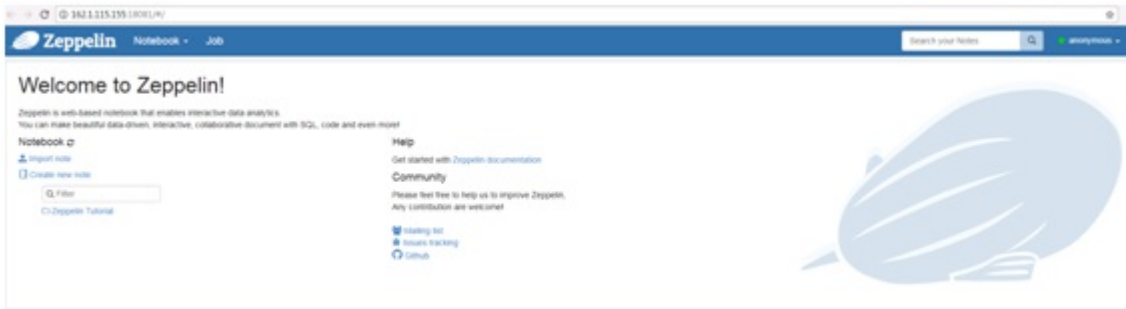
```
sed -i 's/8080/18081/' zeppelin-site.xml
```

```
<property>
  <name>zepelin.server.port</name>
  <value>18081</value>
  <description>Server port.</description>
</property>
```

- 运行zeppelin

```
cd /opt/zeppelin-0.7.2-bin-all/
./bin/zeppelin-daemon.sh start
```

- 在浏览器中输入地址zeppelin_ip:18081登陆， zeppelin_ip为安装zeppelin的虚拟机IP。



- 根据产品文档创建用户test，并赋予足够权限，下载用户test的keytab文件user.keytab，上传至/opt/目录下。
- 编辑zeppelin-site.xml文件，将zeppelin.anonymous.allowed参数的true修改为false。

```
<property>
  <name>zepelin.anonymous.allowed</name>
  <value>false</value>
  <description>Anonymous user allowed by default</description>
</property>
```

- 编辑shiro.ini文件，位置/opt/zeppelin-0.7.2-bin-all/conf/shiro.ini

```
cp shiro.ini.template shiro.ini
vi shiro.ini
```

[urls]authc表示对任何url访问都需要验证

```
[urls]
# This section is used for url-based security.
# You can secure interpreter, configuration and credential information by :
# anon means the access is anonymous.
# authc means Form based Auth Security
# To enforce security, comment the line below and uncomment the next one
/api/version = anon
#/api/interpreter/** = authc, roles[admin]
#/api/configurations/** = authc, roles[admin]
#/api/credential/** = authc, roles[admin]
#/** = anon
/** = authc
```

[users]下增加用户test，密码Huawei@123

```
[users]
# List of users with their password allowed to access Zeppelin.
# To use a different strategy (LDAP / Database / ...) check the shi
admin = password1, admin
user1 = password2, role1, role2
user2 = password3, role3
user3 = password4, role2
test = Huawei@123
```

- 重启zeppelin。

```
cd /opt/zeppelin-0.7.2-bin-all/
./bin/zeppelin-daemon.sh restart
```

- 使用test用户登陆Zeppelin

Zeppelin连接Hive

操作场景

Zeppelin中配置JDBC解析器，对接Hive的JDBC接口。

前提条件

- 已经完成Zeppelin 0.7.2的安装；
- 已完成FusionInsight HD客户端的安装，包含Hive组件。

操作步骤

- 将 `/opt/hadoopclient/Hive/Beeline/lib/` 下的jar包拷贝至 `/opt/zeppelin-0.7.2-bin-all/interpreter/jdbc/` 目录下。
- 将从新拷贝过来的jar包的属主和权限修改为和`opt/zeppelin-0.7.2-bin-all/interpreter/jdbc/`下原有的jar包相同

```
chown 501:wheel *.jar
chmod 644 *.jar
```

- 编辑zeppelin-env.sh文件，位置`opt/zeppelin-0.7.2-bin-all/conf`，加入以下三个配置内容

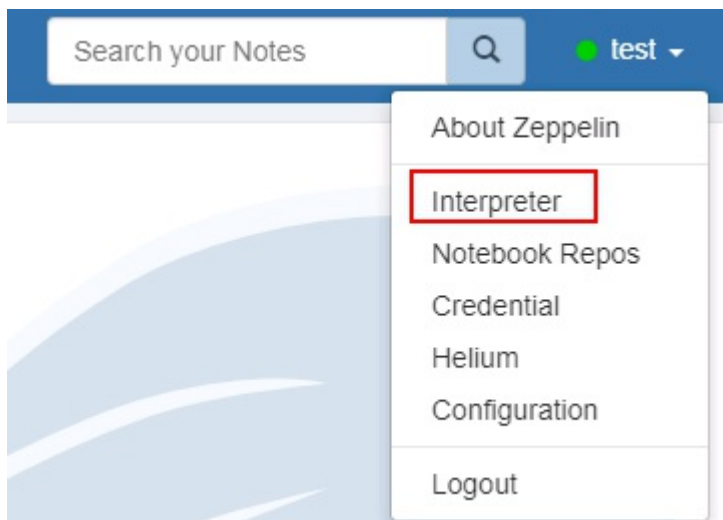
```
export JAVA_HOME=/opt/hadoopclient/JDK/jdk
export ZEPPELIN_INTP_JAVA_OPTS="-Djava.security.krb5.conf=/etc/krb5.conf -Djava.security.auth.login.config=/opt/zeppelin-0.7.2-bin-all/conf/jaas.conf -Dzookeeper.server.principal=zookeeper/hadoop.hadoop.com -Dzookeeper.request.timeout=120000"
export HADOOP_CONF_DIR=/opt/hadoopclient/HDFS/hadoop/etc/hadoop
```

- 从FusionInsight客户端下载用户test的userkeytab和krb5.conf文件，将krb5.conf文件放在/etc/下
- 使用 `vi /opt/zeppelin-0.7.2-bin-all/conf/` 新建hbase的认证文件jaas.conf，内容如下：

```
Client {
com.sun.security.auth.module.Krb5LoginModule required
useKeyTab=true
keyTab="/opt/user.keytab"
principal="test"
useTicketCache=false
storeKey=true
debug=true;
};
```

其中用户为在FusionInsight Manager中创建的test用户，将test的keytab文件user.key放在/opt/目录下

- 登陆Zeppelin，选择右上角菜单中的 Interpreter



- 选择JDBC，点击 **edit** 编辑，修改default.driver和default.url参数，点击 **save** 保存

```
default.driver: org.apache.hive.jdbc.HiveDriver
```

```
default.url:
jdbc:hive2://162.1.93.103:24002,162.1.93.102:24002,162.1.93.101:24002/;serviceDiscoveryMode=zooKeeper;zooKeeperNamespace=hiveserver2;sasl.qop=auth-conf;auth=KERBEROS;principal=hive/hadoop.hadoop.com@HADOOP.COM;user.principal=test;user.keytab=/opt/user.keytab
```

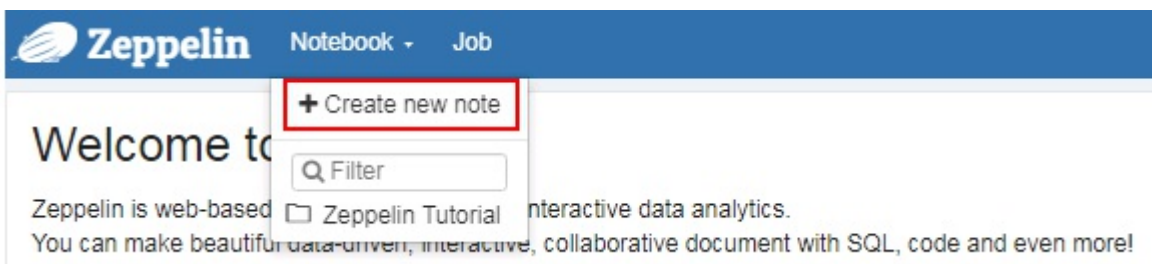
Properties

name	value
common.max_count	1000
default.driver	org.apache.hive.jdbc.HiveDriver
default.password	
default.url	jdbc:hive2://162.1.93.103:24002,162.1.93.102:24002,162.1.93.101:24002/;serviceDiscoveryMode=zooKeeper;zooKeeperNamespace=hiveserver2;sasl.qop=auth-conf;auth=KERBEROS;principal=hive/hadoop.hadoop.com@HADOOP.COM;user.principal=test;user.keytab=/opt/user.keytab
default.user	gpadmin
zeppelin.interpreter.localRepo	/opt/zeppelin-0.7.2-bin-all/local-repo/2CPF6R51K

- 重启zeppelin。

```
source /opt/hadoopclient/bigdata_env
kinit -kt /opt/user.keytab test
cd /opt/zeppelin-0.7.2-bin-all/bin
./zeppelin-daemon.sh restart
```

- 页面选择Notebook -> Create new note



- 自定义note名称，例如hive

Create new note

Note Name

hive 自定义名称

Default Interpreter jdbc

Use '/' to create folders. Example: /NoteDirA/Note1

Create Note

- 编辑note，点击右侧“执行”按钮。

```
%jdbc
Show tables;
Select * from workers_info;
```

- 查看结果



Zeppelin连接HBase

操作场景

Zeppelin中配置Hbase解析器，对接Hbase

前提条件

- 已经完成Zeppelin 0.7.2的安装；
- 已完成FusionInsight HD客户端的安装，包含HBase组件。

操作步骤

- 将 `/opt/hadoopclient/HBase/hbase/lib/` 以下的jar包拷贝至 `/opt/zeppelin-0.7.2-bin-all/interpreter/hbase/` 目录下，overwrite选择n

```
[root@localhost hbase]# cp /opt/hadoopclient/HBase/hbase/lib/*.jar .  
cp: overwrite './activation-1.1.jar'? n  
cp: overwrite './aopalliance-1.0.jar'? n  
cp: overwrite './apacheds-i18n-2.0.0-M15.jar'? n  
cp: overwrite './apacheds-kerberos-codec-2.0.0-M15.jar'? n  
cp: overwrite './api-asn1-api-1.0.0-M20.jar'? n  
cp: overwrite './api-util-1.0.0-M20.jar'? n  
cp: overwrite './asm-3.1.jar'? n  
cp: overwrite './avro-1.7.4.jar'? n
```

- 在/opt/zeppelin-0.7.2-bin-all/interpreter/hbase/下新建目录zeppelin_hbase_jar

```
mkdir /opt/zeppelin-0.7.2-bin-all/interpreter/hbase/zeppelin_hbase_jar
```

- 将/opt/zeppelin-0.7.2-bin-all/interpreter/hbase/下与FusionInsight冲突的38个jar包移动到zeppelin_hbase_jar目录中
 - o commons-codec-1.5.jar
 - o commons-collections-3.2.1.jar
 - o commons-configuration-1.9.jar
 - o commons-lang-2.5.jar
 - o commons-logging-1.1.1.jar
 - o guava-15.0.jar
 - o hadoop-annotations-2.6.0.jar
 - o hadoop-auth-2.5.1.jar
 - o hadoop-client-2.5.1.jar
 - o hadoop-common-2.5.1.jar
 - o hadoop-hdfs-2.5.1.jar
 - o hadoop-mapreduce-client-app-2.5.1.jar
 - o hadoop-mapreduce-client-common-2.5.1.jar
 - o hadoop-mapreduce-client-core-2.5.1.jar
 - o hadoop-mapreduce-client-jobclient-2.5.1.jar
 - o hadoop-mapreduce-client-shuffle-2.5.1.jar
 - o hadoop-yam-api-2.6.0.jar
 - o hadoop-yam-client-2.5.1.jar
 - o hadoop-yam-common-2.6.0.jar
 - o hadoop-yam-server-common-2.5.1.jar

- hbase-annotations-1.0.0.jar
 - hbase-client-1.0.0.jar
 - hbase-common-1.0.0.jar
 - hbase-common-1.0.0-tests.jar
 - hbase-hadoop2-compat-1.0.0.jar
 - hbase-hadoop-compat-1.0.0.jar
 - hbase-prefix-tree-1.0.0.jar
 - hbase-protocol-1.0.0.jar
 - hbase-server-1.0.0.jar
 - httpClient-4.5.1.jar
 - httpcore-4.4.1.jar
 - jettison-1.1.jar
 - netty-3.6.2.Final.jar
 - slf4j-api-1.7.10.jar
 - slf4j-log4j12-1.7.10.jar
 - xmlenc-0.52.jar
 - zookeeper-3.4.6.jar
- 最终/opt/zeppelin-0.7.2-bin-all/interpreter/hbase/有152个jar包

```
[root@localhost hbase]# cd zeppelin_hbase_jar/
[root@localhost zeppelin_hbase_jar]# ll |wc -l
38
[root@localhost zeppelin_hbase_jar]# cd ../
[root@localhost hbase]# ll | wc -l
152
[root@localhost hbase]#
```

- 编辑zeppelin-env.sh文件，位置/opt/zeppelin-0.7.2-bin-all/conf，加入以下三个配置内容

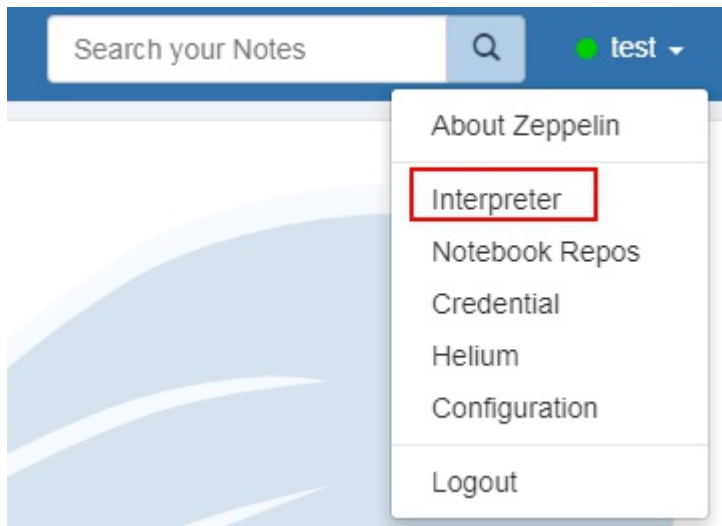
```
export JAVA_HOME=/opt/hadoopclient/JDK/jdk
export ZEPPELIN_INTP_JAVA_OPTS="-Djava.security.krb5.conf=/etc/krb5.conf -Djava.security.auth.login.config=/opt/zeppelin-0.7.2-bin-all/conf/jaas.conf -Dzookeeper.server.principal=zookeeper/hadoop.hadoop.com -Dzookeeper.request.timeout=120000"
export HBASE_HOME=/opt/hadoopclient/HBase/hbase
```

- 从FusionInsight客户端下载用户test的user.keytab和krb5.conf文件，将krb5.conf文件放在/etc/下
- 使用 `vi /opt/zeppelin-0.7.2-bin-all/conf/` 新建hbase的认证文件jaas.conf，内容如下：

```
Client {
com.sun.security.auth.module.Krb5LoginModule required
useKeyTab=true
keyTab="/opt/user.keytab"
principal="test"
useTicketCache=false
storeKey=true
debug=true;
};
```

其中用户为在FusionInsight Manager中创建的test用户，将test的keytab文件user.key放在/opt/目录下

- 登陆Zeppelin，选择右上角菜单中的 Interpreter



- 选择hbase, 点击 **edit** 编辑, 修改hbase.home参数, 点击 **save** 保存

```
hbase.home: /opt/hadoopclient/HBase/hbase
```


hbase.home

/opt/hadoopclient/HBase/hbase

- 重启zeppelin

```
source /opt/hadoopclient/bigdata_env
kinit -kt /opt/user.keytab test
cd /opt/zeppelin-0.7.2-bin-all/bin
./zeppelin-daemon.sh restart
```

- 页面选择Notebook -> Create new note



- 自定义note名称, 例如hbase

Create new note

Note Name

hbase

Default Interpreter

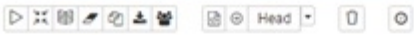
hbase

Use '/' to create folders. Example: /NoteDirA/Note1

Create Note

- 编辑note, 点击右侧 **执行** 按钮

```
%hbase
create 'test2', 'cf'
put 'test2', 'row1', 'cf:a', 'value1'
```

```
hbase 
$sh
source /opt/hadoopclient/bigdata_env
kinit -kt /opt/user.keytab test
Took 1 sec. Last updated by test at August 11 2017, 4:45:51 PM.

hbase
help

Ruby Hashes. They look like this:
{'key1' => 'value1', 'key2' => 'value2', ...}
and are opened and closed with curly-braces. Key/values are delimited by the
'=' character combination. Usually keys are predefined constants such as
NAME, VERSIONS, COMPRESSION, etc. Constants do not need to be quoted. Type
'Object.constants' to see a (messy) list of all constants in the environment.
If you are using binary keys or values and need to enter them in the shell, use
double-quoted hexadecimal representation. For example:
hbase> get 't1', "key\x03\x3f\xcd"
hbase> get 't1', "key\003\023\011"
hbase> put 't1', "test\xef\xff", 'f1:', "\x01\x33\x40"
The HBase shell is the (J)Ruby IRB with the above HBase-specific commands added.
For more on the HBase Shell, see http://hbase.apache.org/book.html
Took 7 sec. Last updated by test at August 11 2017, 4:45:59 PM.

hbase
create 'test2', 'cf'
put 'test2', 'row1', 'cf:a', 'value1'

0 row(s) in 0.4490 seconds
0 row(s) in 0.0950 seconds
Took 1 sec. Last updated by test at August 11 2017, 4:46:39 PM.
```

- 在FusionInsight的客户端下可以看到创建的hbase表test2和数据

```
hbase(main):003:0> scan "test2"
ROW                                COLUMN+CELL
  row1                             column=cf:a, timestamp=1502441199545, value=value1
1 row(s) in 0.0300 seconds
```

Zeppelin连接Spark

操作场景

Zeppelin中配置Spark解析器

前提条件

- 完成Zeppelin0.7.2的安装;
- 已完成FusionInsight HD V100R002C60U20和客户端的安装, 包含Spark组件。
- 参考<http://zeppelin.apache.org/docs/latest/interpreter/spark.html>

操作步骤

- 将 `/opt/zeppelin-0.7.2-bin-all/lib/` 目录下的原有的相关的jar包删除
 - hadoop-auth-2.6.0.jar
 - hadoop-common-2.6.0.jar
 - scala-compiler-2.11.7.jar
 - scala-library-2.11.7.jar
 - scala-parser-combinators_2.11-1.0.4.jar
 - scala-reflect-2.11.7.jar
 - scala-xml_2.11-1.0.2.jar
- 将 `/opt/hadoopclient/Spark/adaptor/dev_lib/` 下的以下jar包拷贝到 `/opt/zeppelin-0.7.2-bin-all/lib/` 目录下
 - hadoop-auth-2.7.2.jar
 - hadoop-common-2.7.2.jar
 - scala-compiler-2.10.4.jar
 - scala-library-2.10.4.jar
 - scala-reflect-2.10.4.jar
- 将 `/opt/zeppelin-0.7.2-bin-all/lib/` 下的jackson的相关jar包删除
 - jackson-annotations-2.5.0.jar
 - jackson-core-2.5.3.jar
 - jackson-core-asl-1.9.13.jar

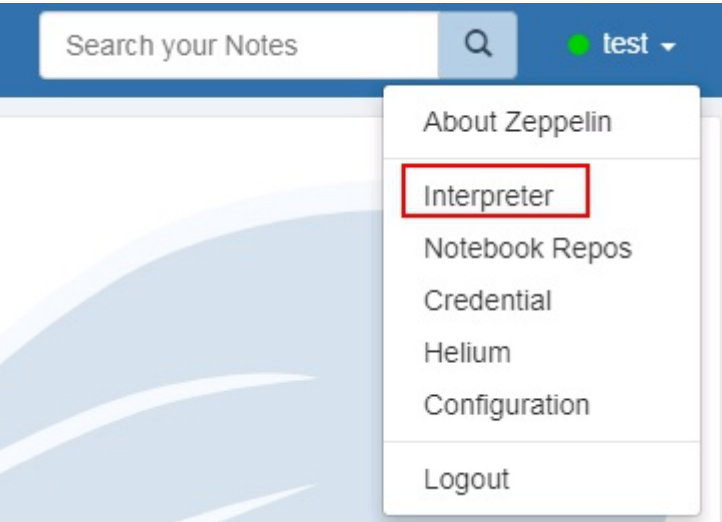
- jackson-databind-2.5.3.jar
- jackson-mapper-asl-1.9.13.jar
- 将 `/opt/hadoopclient/Spark/adapter/dev_lib/` 下的jackson相关的jar包拷贝到 `/opt/zeppelin-0.7.2-bin-all/lib/` 下
 - jackson-annotations-2.4.0.jar
 - jackson-core-2.4.4.jar
 - jackson-core-asl-1.9.13.jar
 - jackson-databind-2.4.4.jar
 - jackson-jaxrs-1.9.13.jar
 - jackson-mapper-asl-1.9.13.jar
 - jackson-module-scala_2.10-2.4.4.jar
 - jackson-xc-1.9.13.jar
- 将步骤1和步骤2所有从spark客户端拷贝过来的jar包的属主和权限修改为和 `/opt/zeppelin-0.7.2-bin-all/lib/` 下原有的jar包相同

```
chown 501:wheel *.jar
chmod 644 *.jar
```

- 编辑zeppelin-env.sh文件，位置 `/opt/zeppelin-0.7.2-bin-all/conf`，加入以下内容

```
export MASTER=yarn-client
export SPARK_HOME=/opt/hadoopclient/Spark/spark
export HADOOP_CONF_DIR=/opt/hadoopclient/HDFS/hadoop/etc/hadoop
```

- 登陆Zeppelin，选择右上角菜单中的 Interpreter



- 选择Spark，点击 **edit** 编辑，将 Master 参数改为 yarn-client，点击 **save** 保存

Properties	
name	value
args	
master	yarn-client
spark.app.name	Zeppelin

- 重启zeppelin

```
source /opt/hadoopclient/bigdata_env
kinit -kt /opt/user.keytab test
cd /opt/zeppelin-0.7.2-bin-all/bin
./zeppelin-daemon.sh restart
```

- 执行zeppelin的spark样例代码zeppelin Tutorial -> Basic Features(Spark)

样例代码需要访问Internet上的资源，所以保证zeppelin所在的节点可以联网，检测是否能打开以下链接

Load data into table

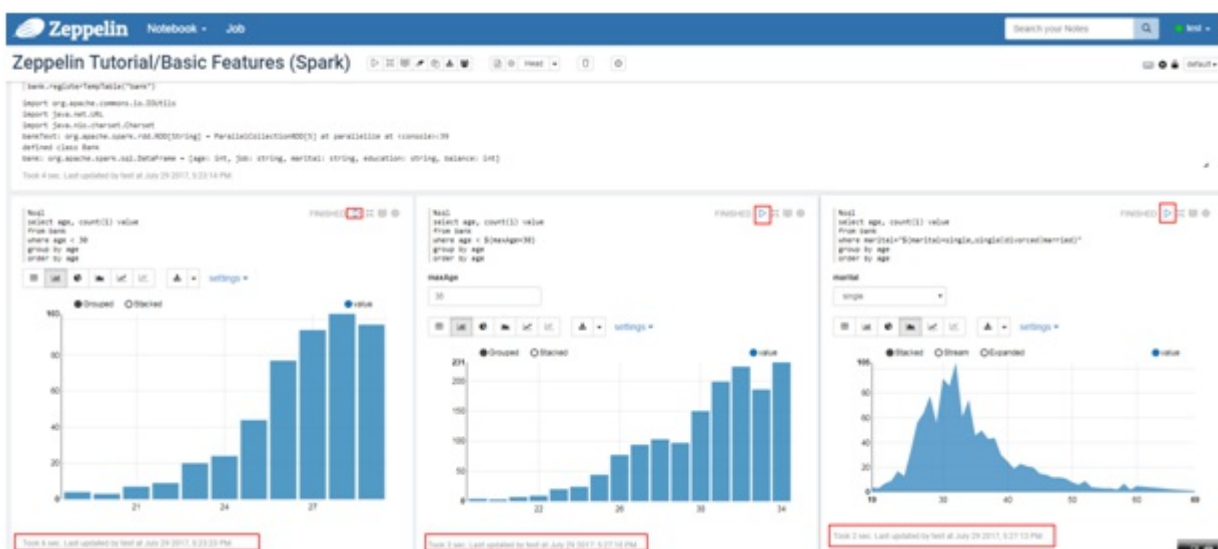
```
import org.apache.commons.io.IOUtils
import java.net.URL
import java.nio.charset.Charset

// Zeppelin creates and injects sc (SparkContext) and sqlContext (HiveContext or SqlContext)
// So you don't need create them manually

// load bank data
val bankText = sc.parallelize(
  IOUtils.toString(
    new URL("https://s3.amazonaws.com/apache-zeppelin/tutorial/bank/bank.csv"),
    Charset.forName("utf8")).split("\n"))

case class Bank(age: Integer, job: String, marital: String, education: String, balance: Integer)

val bank = bankText.map(s => s.split(";")).filter(s => s(0) != "\"age\"").map(
  s => Bank(s(0).toInt,
    s(1).replaceAll("\"", ""),
    s(2).replaceAll("\"", ""),
    s(3).replaceAll("\"", ""),
    s(5).replaceAll("\"", "").toInt)
).toDF()
bank.registerTempTable("bank")
```



- 执行zeppelin的spark样例代码Zeppelin Tutorial -> Matplotlib (Python • PySpark)

安装python-matplotlib

```
yum install python-matplotlib
```

安装Anaconda2-4.4

```
wget https://repo.continuum.io/archive/Anaconda2-4.4.0-Linux-x86_64.sh
sh Anaconda2-4.4.0-Linux-x86_64.sh
```

配置环境变量PATH，将python换成安装Anaconda安装目录中的python

```
export PATH=/root/anaconda2/bin:$PATH
```

在zeppelin的界面中，选择右上角的 Interpreter

选择Spark，点击 **edit** 编辑，将 zeppelin.pyspark.python 参数改为Anaconda安装目录中的python，点击 **save** 保存

zeppelin.pyspark.python


/root/anaconda2/bin/python

执行zeppelin的pyspark样例代码Zeppelin Tutorial -> Matplotlib

First line (figure will be displayed here)

FINISHED ▶

```
%spark
import matplotlib.pyplot as plt
plt.close() # Added here to reset the plot when rerunning the paragraph
%matplotlib inline
plt.figure(figsize=(10, 5))
plt.plot([1, 2, 3], label='y=x')
[<matplotlib.lines.Line2D object at 0x7fa9c369ef00>]
```



Took 30 sec. Last updated by test at July 31 2017, 5:15:10 PM

Zeppelin连接SparkR

操作场景

Zeppelin中配置Spark解析器，连接SparkR

前提条件

- 完成Zeppelin0.7.2的安装;
- 已完成FusionInsight HD V100R002C60U20和客户端的安装，包含Spark组件。
- 参考<http://zeppelin.apache.org/docs/latest/interpreter/spark.html>

操作步骤

- 由于Spark的Executor上也需要执行R，所以除了在Zeppelin的节点上安装R以外，所有FusionInsight集群节点上也要安装同版本的R，安装步骤如下：

不同OS配置yum源时下载的文件路径有所不同，下面以Redhat6.6安装R为例

如果安装R的节点无法访问互联网，参考FAQ进行R的安装

- 配置Redhat6.6的yum源

```
cd ~
rpm -aq | grep yum | xargs rpm -e --nodeps
wget http://mirrors.163.com/centos/6/os/x86_64/Packages/python-iniparse-0.3.1-2.1.el6.noarch.rpm
wget http://mirrors.163.com/centos/6/os/x86_64/Packages/yum-metadata-parser-1.1.2-16.el6.x86_64.rpm
wget http://mirrors.163.com/centos/6/os/x86_64/Packages/yum-3.2.29-81.el6.centos.noarch.rpm
wget http://mirrors.163.com/centos/6/os/x86_64/Packages/yum-plugin-fastestmirror-1.1.30-40.el6.noarch.rpm
wget http://mirrors.163.com/centos/6/os/x86_64/Packages/python-urlgrabber-3.9.1-11.el6.noarch.rpm
rpm -ivh python-iniparse-0.3.1-2.1.el6.noarch.rpm
rpm -ivh yum-metadata-parser-1.1.2-16.el6.x86_64.rpm
rpm -U python-urlgrabber-3.9.1-11.el6.noarch.rpm
rpm -ivh yum-3.2.29-81.el6.centos.noarch.rpm yum-plugin-fastestmirror-1.1.30-40.el6.noarch.rpm
cd /etc/yum/repos.d/
wget http://mirrors.163.com/.help/CentOS6-Base-163.repo
sed -i 's/$releasever/6/g' /etc/yum/repos.d/CentOS6-Base-163.repo
yum clean all
yum makecache
```

- 配置EPEL的源

Redhat 6.x 使用下面命令安装

```
rpm -Uvh https://mirrors.tuna.tsinghua.edu.cn/epel/6/x86_64/epel-release-6-8.noarch.rpm
```

- 更新cache

```
yum clean all
yum makecache
```

- 执行 `yum install R` 安装R的相关的包
- 执行 `R`，检查R是否可用

正常启动如下图所示

```
[root@work opt]# R

R version 3.4.1 (2017-06-30) -- "Single Candle"
Copyright (C) 2017 The R Foundation for Statistical Computing
Platform: x86_64-pc-linux-gnu (64-bit)

R is free software and comes with ABSOLUTELY NO WARRANTY.
You are welcome to redistribute it under certain conditions.
Type 'license()' or 'licence()' for distribution details.

  Natural language support but running in an English locale

R is a collaborative project with many contributors.
Type 'contributors()' for more information and
'citation()' on how to cite R or R packages in publications.


Type 'demo()' for some demos, 'help()' for on-line help, or
'help.start()' for an HTML browser interface to help.
Type 'q()' to quit R.

> 
```

- FusionInsight客户端下测试是否可以使用sparkR

```
source /opt/hadoopclient/bigdata_env
kinit test
sparkR
```

- 正常启动如下图所示

 **Zeppelin** Notebook ▾ Job

Zeppelin Tutorial/R (SparkR)

▶

⌂

📄

✍

📁

📥

👤

📄

🔄

Head ▾

🗑

🕒

Hello R

```
%r
foo <- TRUE
print(foo)
bare <- c(1, 2.5, 4)
print(bare)
double <- 15.0
print(double)
```

[1] TRUE

[1] 1.0 2.5 4.0

[1] 15

Took 0 sec. Last updated by test at August 07 2017, 4:32:26 PM.

FAQ

• FusionInsight集群不允许访问网络，如何安装R

在集群外同版本的Redhat版本下按照本文中yum源的方式进行安装R的操作，最后一步不要执行 `yum install R`

执行 `yum install yum-utils` 安装yumdownloader

执行 `yumdownloader R --resolve --destdir=/tmp/packages` 把所有的rpm安装包下载到 `/tmp/packages` 中

将 `/tmp/packages` 中的所有rpm包复制到集群每个节点的 `/tmp/packages` 中

切换到集群每个节点的 `/tmp/packages` 中，执行 `yum localinstall *.rpm` 完成安装

• 连接hbase出现AuthFialed for /hwbakup/hbase

```
%hbase
help

org.apache.zookeeper KeeperException$AuthFailedException: KeeperErrorCode = AuthFailed for /hwbakup/hbase
    at org.apache.zookeeper KeeperException.create(KeeperException.java:123)
    at org.apache.zookeeper KeeperException.create(KeeperException.java:51)
    at org.apache.zookeeper ZooKeeper.create(ZooKeeper.java:783)
    at org.apache.hadoop.hbase.zookeeper.RecoverableZooKeeper.createNonSequential(RecoverableZooKeeper.java:512)
    at org.apache.hadoop.hbase.zookeeper.RecoverableZooKeeper.create(RecoverableZooKeeper.java:491)
    at org.apache.hadoop.hbase.zookeeper.ZKUtil.createWithParents(ZKUtil.java:1252)
    at org.apache.hadoop.hbase.zookeeper.ZKUtil.createWithParents(ZKUtil.java:1230)
    at com.huawei.hadoop.hbase.backup.zookeeper.BackupZooKeeperWatcher.createBaseZNodes(BackupZooKeeperWatcher.java:137)
    at com.huawei.hadoop.hbase.backup.zookeeper.BackupZooKeeperWatcher.<init>(BackupZooKeeperWatcher.java:78)
    at com.huawei.hadoop.hbase.backup.client.BackupAdmin.<init>(BackupAdmin.java:98)
    at sun.reflect.NativeConstructorAccessorImpl.newInstance0(Native Method)
    at sun.reflect.NativeConstructorAccessorImpl.newInstance(NativeConstructorAccessorImpl.java:57)
    at sun.reflect.DelegatingConstructorAccessorImpl.newInstance(DelegatingConstructorAccessorImpl.java:45)
    at java.lang.reflect.Constructor.newInstance(Constructor.java:526)
    at org.jruby.javasupport.JavaConstructor.newInstanceDirect(JavaConstructor.java:275)
```

原因: zeppelin的原理hbase的jar包与从FusionInsight客户端下拷贝过来的jar冲突。

解决: 将zeppelin中原有的重名jar包移走或删除，全部用FusionInsight客户端下的相关jar包。

• Zeppelin连接spark是报如下 NoSuchMethodError

```
java.lang.NoSuchMethodError: scala.reflect.api.JavaUniverse.runtimeMirror(Ljava/lang/ClassLoader;)Lscala/reflect/api/JavaMirrors$JavaMirror;
    at org.apache.spark.repl.SparkILoop.<init>(SparkILoop.scala:936)
    at org.apache.spark.repl.SparkILoop.<init>(SparkILoop.scala:70)
    at org.apache.zeppelin.spark.SparkInterpreter.open(SparkInterpreter.java:775)
    at org.apache.zeppelin.interpreter.LazyOpenInterpreter.open(LazyOpenInterpreter.java:70)
    at org.apache.zeppelin.spark.PySparkInterpreter.getSparkInterpreter(PySparkInterpreter.java:564)
    at org.apache.zeppelin.spark.PySparkInterpreter.createGatewayServerAndStartScript(PySparkInterpreter.java:208)
    at org.apache.zeppelin.spark.PySparkInterpreter.open(PySparkInterpreter.java:162)
    at org.apache.zeppelin.interpreter.LazyOpenInterpreter.open(LazyOpenInterpreter.java:70)
    at org.apache.zeppelin.interpreter.remote.RemoteInterpreterServer$InterpreterJob.jobRun(RemoteInterpreterServer.java:491)
    at org.apache.zeppelin.scheduler.Job.run(Job.java:175)
    at org.apache.zeppelin.scheduler.FIFO Scheduler$1.run(FIFO Scheduler.java:139)
    at java.util.concurrent.Executors$RunnableAdapter.call(Executors.java:471)
    at java.util.concurrent.FutureTask.run(FutureTask.java:262)
    at java.util.concurrent.ScheduledThreadPoolExecutor$ScheduledFutureTask.access$201(ScheduledThreadPoolExecutor.java:178)
    at java.util.concurrent.ScheduledThreadPoolExecutor$ScheduledFutureTask.run(ScheduledThreadPoolExecutor.java:292)
    at java.util.concurrent.ThreadPoolExecutor.runWorker(ThreadPoolExecutor.java:1145)
    at java.util.concurrent.ThreadPoolExecutor$Worker.run(ThreadPoolExecutor.java:615)
```

原因: jar包冲突

解决: 删除 `/opt/zeppelin-0.7.2-bin-all/lib/` 下原有jar包scala-reflect-2.11.7.jar，替换为FusionInsight客户端下的jar包，重启zeppelin

- Zeppelin执行Spark样例代码时报GC overhead limit exceeded

```
Load data into table

import org.apache.commons.io.IOUtils
import java.net.URL
import java.nio.charset.Charset

// Zeppelin creates and injects sc (SparkContext) and sqlContext (HiveContext or SqlContext)
// So you don't need create them manually

// load bank data
val bankText = sc.parallelize(
  IOUtils.toString(
    new URL("https://s3.amazonaws.com/apache-zeppelin/tutorial/bank/bank.csv"),
    Charset.forName("utf8")).split("\n"))

case class Bank(age: Integer, job: String, marital: String, education: String, balance: Integer)

val bank = bankText.map(s => s.split(";")).filter(s => s(0) != "\age\").map(
  s => Bank(s(0).toInt,
    s(1).replaceAll("\", ""),
    s(2).replaceAll("\", ""),
    s(3).replaceAll("\", ""),
    s(5).replaceAll("\", ").toInt
  )
).toDF()
bank.registerTempTable("bank")

java.lang.OutOfMemoryError: GC overhead limit exceeded
```

原因：内存不够

解决：安装Zeppelin的节点的内存需要16G以上

- 执行zeppelin的样例代码Zeppelin Tutorial/Matplotlib (Python PySpark)报如下错误

```
Add title ERROR ▶ ✖ ⓘ ⚙

%pyspark
plt.title('Inline plotting example', fontsize=20)

<matplotlib.text.Text object at 0x2af1d90>

Traceback (most recent call last):
  File "/tmp/zeppelin_pyspark-6355775624574224283.py", line 367, in <module>
    raise Exception(traceback.format_exc())
Exception: Traceback (most recent call last):
  File "/tmp/zeppelin_pyspark-6355775624574224283.py", line 365, in <module>
    exec(code, __zcUserQueryNamespace)
  File "<stdin>", line 2, in <module>
  File "/opt/zeppelin-0.7.2-bin-all/interpreter/lib/python/backend_zinline.py", line 303, in displayhook
    show()
  File "/opt/zeppelin-0.7.2-bin-all/interpreter/lib/python/backend_zinline.py", line 72, in __call__
    manager.show(**kwargs)
  File "/opt/zeppelin-0.7.2-bin-all/interpreter/lib/python/backend_zinline.py", line 210, in show
    self.canvas.draw_idle()
  File "/opt/zeppelin-0.7.2-bin-all/interpreter/lib/python/backend_zinline.py", line 134, in draw_idle
    if not self._is_idle_drawing:
AttributeError: 'FigureCanvasZInline' object has no attribute '_is_idle_drawing'

Took 0 sec. Last updated by test at August 01 2017, 8:17:53 PM.
```

原因：python版本问题

解决：安装Anaconda2-4.4