

Zeppelin对接FusionInsight HD

适用场景

Zeppelin 0.7.3 ↔ FusionInsight HD V100R002C70SPC100 (Spark2.x)

编译Zeppelin

- 安装maven:

```
wget http://apache.osuosl.org/maven/maven-3/3.3.9/binaries/apache-maven-3.3.9-bin.tar.gz
tar -xzf apache-maven-3.3.9-bin.tar.gz -C /opt/
```

修改profile文件 `vi /etc/profile` ,增加以下配置

```
export PATH=$PATH:/opt/apache-maven-3.3.9/bin
```

导入环境变量

```
source /etc/profile
```

执行 `mvn -v`

- 安装git

```
yum install -y git
```

- 安装nodejs:

```
wget https://nodejs.org/dist/v6.10.0/node-v6.10.0-linux-x64.tar.xz --no-check-certificate
tar -xvf node-v6.10.0-linux-x64.tar.xz -C /opt/
```

修改profile文件 `vi /etc/profile` ,增加以下配置

```
export PATH=$PATH:/opt/apache-maven-3.3.9/bin:/opt/node-v6.10.0-linux-x64/bin
```

导入环境变量

```
source /etc/profile
```

执行 `npm -v`

- 安装bower

```
npm install -g bower
```

配置bower允许root用户执行

```
echo '{ "allow_root": true }' > /root/.bowerrc
```

执行 `bower -v`

- 获取Zeppelin0.7.3的版本

```
git clone https://github.com/apache/zeppelin.git
cd zeppelin
git checkout v0.7.3
```

- 修改scala版本，适配FusionInsight_HD_V100R002C70SPC100的Hadoop版本

在zeppelin代码根目录执行 `vi ./dev/change_scala_version.sh`，修改下图的SCALA_LIB_VERSION为2.11.8

```
check_scala_version "${TO_VERSION}"

if [ "${TO_VERSION}" = "2.11" ]; then
  FROM_VERSION="2.10"
  SCALA_LIB_VERSION="2.11.8"
else
  FROM_VERSION="2.11"
  SCALA_LIB_VERSION="2.10.5"
fi
```

执行命令完成scala版本的修改

```
./dev/change_scala_version.sh 2.11
```

- 执行 `vi pom.xml` 文件的修改<libthrift.version>为0.9.3

```
<!-- common library versions -->
<slf4j.version>1.7.10</slf4j.version>
<log4j.version>1.2.17</log4j.version>
<libthrift.version>0.9.3</libthrift.version>
<gson.version>2.2</gson.version>
<guava.version>15.0</guava.version>
<jetty.version>9.2.15.v20160210</jetty.version>
```

- 执行 `vi hbase/pom.xml` 修改hbase版本和hadoop版本

```
<properties>
  <!--library versions-->
  <hbase.hbase.version>1.0.2</hbase.hbase.version>
  <hbase.hadoop.version>2.7.2</hbase.hadoop.version>
  <jruby.version>1.6.8</jruby.version>
  <protobuf.version>2.5.0</protobuf.version>
  <commons.exec.version>1.1</commons.exec.version>
  <jline.version>2.12.1</jline.version>
```

- 编译Zeppelin

```
mvn clean package -Pbuild-distr -Pspark-2.1 -Dspark.version=2.1.0 -Dhadoop.version=2.7.2 -Phadoop-2.7 -Pscala-2.11 -Psparkr -
DskipTests
```

- 编译完成后在 `zeppelin-distribution/target` 目录下生成 `zeppelin-0.7.3.tar.gz` 文件

安装Zeppelin

操作场景

安装Zeppelin0.7.3

前提条件

- 已完成FusionInsight HD客户端的安装。

操作步骤

- 将编译好的zeppelin-0.7.3.targz上传放到/opt目录下，解压生成zeppelin-0.7.3目录。

```
cp zeppelin-distribution/target/zeppelin-0.7.3.tar.gz /opt
cd /opt
tar -zxvf zeppelin-0.7.3.tar.gz
```

- 配置Zeppelin环境变量，在profile文件中加入如下变量

```
vi /etc/profile
export ZEPPELIN_HOME=/opt/zeppelin-0.7.3
export PATH=$ZEPPELIN_HOME/bin:$PATH
```

- 编辑zeppelin-env.sh文件，位置/opt/zeppelin-0.7.3/conf

```
cd /opt/zeppelin-0.7.3/conf/
cp zeppelin-env.sh.template zeppelin-env.sh
vi zeppelin-env.sh
```

加入如下内容：

```
export JAVA_HOME=/opt/hadoopclient/JDK/jdk
```

编辑zeppelin-site.xml文件，位置/opt/zeppelin-0.7.3/conf/

```
cp zeppelin-site.xml.template zeppelin-site.xml
```

将zeppelin-site.xml中端口8080替换成18081（可自定义，也可以不改）

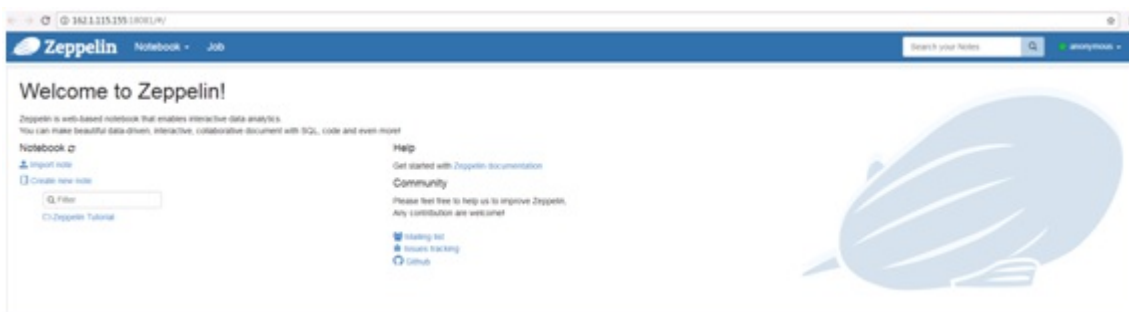
```
sed -i 's/8080/18081/' zeppelin-site.xml
```

```
<property>
  <name>zeppelin.server.port</name>
  <value>18081</value>
  <description>Server port.</description>
</property>
```

- 运行zeppelin

```
cd /opt/zeppelin-0.7.3/
./bin/zeppelin-daemon.sh start
```

- 在浏览器中输入地址zeppelin_ip:18081登陆，zeppelin_ip为安装zeppelin的虚拟机IP。



- 根据产品文档创建用户test，并赋予足够权限，下载用户test的keytab文件user.keytab，上传至/opt/目录下。
- 编辑zeppelin-site.xml文件，将zeppelin.anonymous.allowed参数的true修改为false。

```
<property>
  <name>zeppelin.anonymous.allowed</name>
  <value>false</value>
  <description>Anonymous user allowed by default</description>
</property>
```

- 编辑shiro.ini文件，位置/opt/zeppelin-0.7.3/conf/shiro.ini

```
cp shiro.ini.template shiro.ini
vi shiro.ini
```

[urls]authc表示对任何url访问都需要验证

```
[urls]
# This section is used for url-based security.
# You can secure interpreter, configuration and credential information by :
# anon means the access is anonymous.
# authc means Form based Auth Security
# To enforce security, comment the line below and uncomment the next one
/api/version = anon
#/api/interpreter/** = authc, roles[admin]
#/api/configurations/** = authc, roles[admin]
#/api/credential/** = authc, roles[admin]
#/** = anon
/** = authc
```

[users]下增加用户test，密码Huawei@123

```
[users]
# List of users with their password allowed to access Zeppelin.
# To use a different strategy (LDAP / Database / ...) check the shi
admin = password1, admin
user1 = password2, role1, role2
user2 = password3, role3
user3 = password4, role2
test = Huawei@123
```

- 重启zeppelin。

```
cd /opt/zeppelin-0.7.3/
./bin/zeppelin-daemon.sh restart
```

- 使用test用户登陆Zeppelin

Zeppelin连接Hive

操作场景

Zeppelin中配置JDBC解析器，对接Hive的JDBC接口。

前提条件

- 已经完成Zeppelin 0.7.3的安装；
- 已完成FusionInsight HD客户端的安装，包含Hive组件。

操作步骤

- 将 /opt/hadoopclient/Hive/Beeline/lib/ 下的jar包拷贝至 /opt/zeppelin-0.7.3/interpreter/jdbc/ 目录下。
- 将从新拷贝过来的jar包的属主和权限修改为和/opt/zeppelin-0.7.3/ interpreter/jdbc/下原有的jar包相同

```
chown 501:wheel *.jar
chmod 644 *.jar
```

- 编辑zeppelin-env.sh文件，位置/opt/zeppelin-0.7.3/conf，加入以下三个配置内容

```
export JAVA_HOME=/opt/hadoopclient/JDK/jdk
export ZEPPELIN_INT_P_JAVA_OPTS="-Djava.security.krb5.conf=/etc/krb5.conf -Djava.security.auth.login.config=/opt/zeppelin-0.7.3/conf/jaas.conf -Dzookeeper.server.principal=zookeeper/hadoop.hadoop.com -Dzookeeper.request.timeout=120000"
export HADOOP_CONF_DIR=/opt/hadoopclient/HDFS/hadoop/etc/hadoop
```

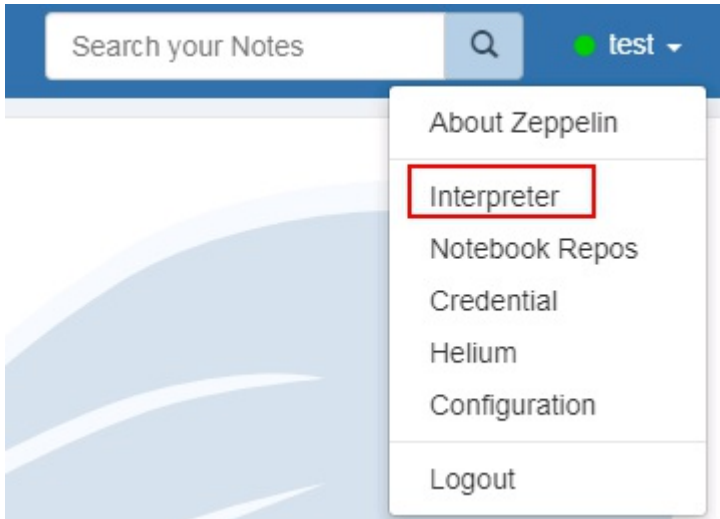
- 从FusionInsight客户端下载用户test的user.keytab和krb5.conf文件，将krb5.conf文件放在/etc/下
- 使用 vi /opt/zeppelin-0.7.3/conf/ 新建hbase的认证文件jaas.conf，内容如下：

```
Client {
com.sun.security.auth.module.Krb5LoginModule required
```

```
useKeyTab=true
keyTab="/opt/user.keytab"
principal="test"
useTicketCache=false
storeKey=true
debug=true;
};
```

其中用户为在FusionInsight Manager中创建的test用户，将test的keytab文件user.key放在/opt/目录下

- 登陆Zeppelin，选择右上角菜单中的 Interpreter



- 选择JDBC，点击 **edit** 编辑，修改default.driver和default.url参数，点击 **save** 保存

```
default.driver: org.apache.hive.jdbc.HiveDriver
default.url:
jdbc:hive2://162.1.93.103:24002,162.1.93.102:24002,162.1.93.101:24002/;serviceDiscoveryMode=zooKeeper;zooKeeperNamespace=hiveserver2;sasl.qop=auth-conf;auth=KERBEROS;principal=hive/hadoop.hadoop.com@HADOOP.COM;user.principal=test;user.keytab=/opt/user.keytab
```


Properties

name	value
common.max_count	1000
default.driver	org.apache.hive.jdbc.HiveDriver
default.password	
default.url	jdbc:hive2://162.1.93.103:24002,162.1.93.102:24002,162.1.93.101:24002/;serviceDiscoveryMode=zooKeeper;zooKeeperNamespace=hiveserver2;sasl.qop=auth-conf;auth=KERBEROS;principal=hive/hadoop.hadoop.com@HADOOP.COM;user.principal=test;user.keytab=/opt/user.keytab
default.user	gpadmin
zeppelin.interpreter.localRepo	/opt/zeppelin-0.7.2-bin-all/local-repo/2CPF6R51K

- 重启zeppelin。

```
source /opt/hadoopclient/bigdata_env
kinit -kt /opt/user.keytab test
cd /opt/zeppelin-0.7.3/bin
./zeppelin-daemon.sh restart
```

- 页面选择Notebook -> Create new note



- Create Note

```
cd /opt/zeppelin-0.7.3/interpreter/hbase
mkdir hbase_jar
mv hbase*.jar hbase_jar
mv hadoop*.jar hbase_jar
mv zookeeper-3.4.6.jar hbase_jar
```

- 将 `/opt/hadoopclient/HBase/hbase/lib/` 以下的jar包拷贝至 `/opt/zeppelin-0.7.3/interpreter/hbase/` 目录下

```
cp /opt/hadoopclient/HBase/hbase/lib/hbase-*.jar /opt/zeppelin-0.7.3/interpreter/hbase
cp /opt/hadoopclient/HBase/hbase/lib/hadoop-*.jar /opt/zeppelin-0.7.3/interpreter/hbase
cp /opt/hadoopclient/HBase/hbase/lib/zookeeper-*.jar /opt/zeppelin-0.7.3/interpreter/hbase
cp /opt/hadoopclient/HBase/hbase/lib/dynallogger-V100R002C30.jar /opt/zeppelin-0.7.3/interpreter/hbase
```

- 编辑zeppelin-env.sh文件，位置`/opt/zeppelin-0.7.3/conf`，加入以下三个配置内容

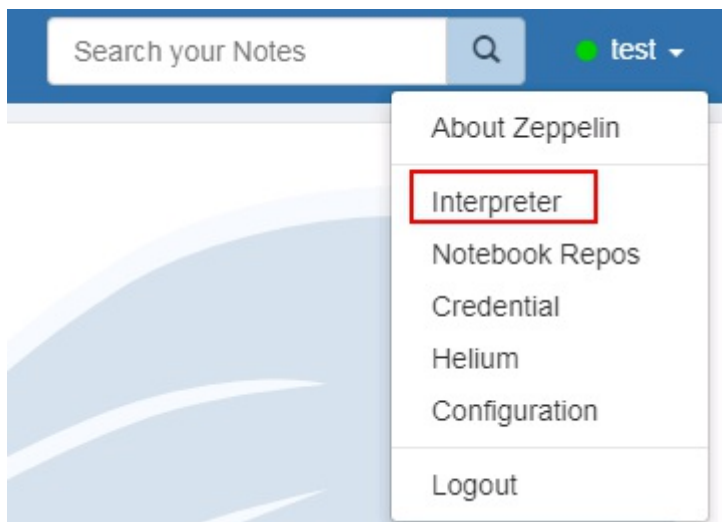
```
export JAVA_HOME=/opt/hadoopclient/JDK/jdk
export ZEPPELIN_INTP_JAVA_OPTS="-Djava.security.krb5.conf=/etc/krb5.conf -Djava.security.auth.login.config=/opt/zeppelin-0.7.3/conf/jaas.conf -Dzookeeper.server.principal=zookeeper/hadoop.hadoop.com -Dzookeeper.request.timeout=120000"
export HBASE_HOME=/opt/hadoopclient/HBase/hbase
```

- 从FusionInsight客户端下载用户test的userkeytab和krb5.conf文件，将krb5.conf文件放在/etc/下
- 使用 `vi /opt/zeppelin-0.7.3/conf/` 新建hbase的认证文件jaas.conf，内容如下：

```
Client {
com.sun.security.auth.module.Krb5LoginModule required
useKeyTab=true
keyTab="/opt/user.keytab"
principal="test"
useTicketCache=false
storeKey=true
debug=true;
};
```

其中用户为在FusionInsight Manager中创建的test用户，将test的keytab文件user.key放在/opt/目录下

- 登陆Zeppelin，选择右上角菜单中的 Interpreter



- 选择hbase，点击 **edit** 编辑，修改hbase.home参数，点击 **save** 保存

```
hbase.home: /opt/hadoopclient/HBase/hbase
```

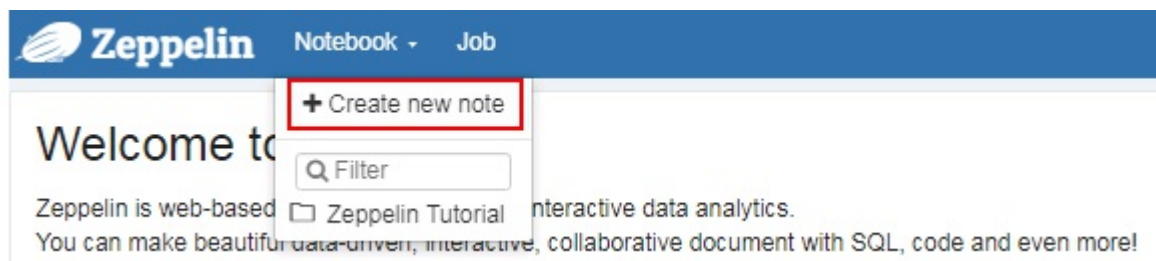
hbase.home

/opt/hadoopclient/HBase/hbase

- 重启zeppelin

```
source /opt/hadoopclient/bigdata_env
kinit -kt /opt/user.keytab test
cd /opt/zeppelin-0.7.3/bin
./zeppelin-daemon.sh restart
```

- 页面选择Notebook -> Create new note



- 自定义note名称，例如hbase

Create new note

Note Name

hbase

Default Interpreter

hbase

Use '/' to create folders. Example: /NoteDirA/Note1

Create Note

- 编辑note，点击右侧 执行 按钮

```
%hbase
create 'test2', 'cf'
put 'test2', 'row1', 'cf:a', 'value1'
```

hbase

Run

```
%hbase
create 'test2', 'cf'
put 'test2', 'row1', 'cf:a', 'value1'
```

0 row(s) in 0.4490 seconds
0 row(s) in 0.0950 seconds
Took 1 sec. Last updated by test at August 11 2017, 4:46:39 PM.

- 在FusionInsight的客户端下可以看到创建的hbase表test2和数据

```
hbase(main):003:0> scan "test2"
ROW COLUMN+CELL
row1 column=cf:a, timestamp=1502441199545, value=value1
1 row(s) in 0.0300 seconds
```

Zeppelin连接Spark

操作场景

前提条件

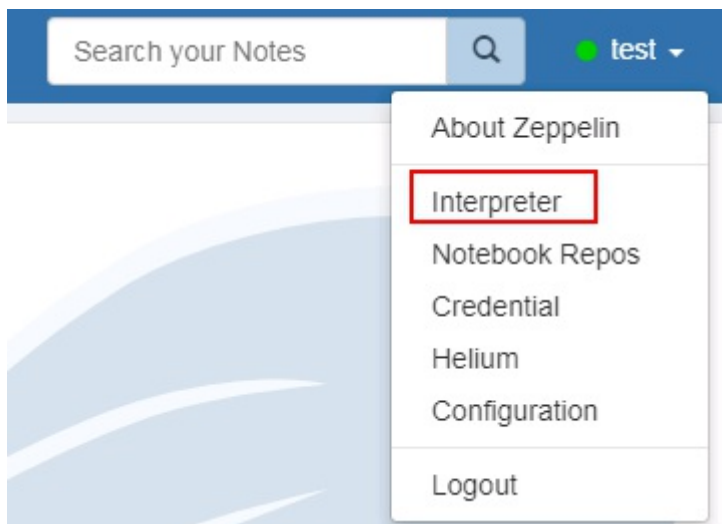
- 完成Zeppelin0.7.3的安装;
- 已完成FusionInsight HD V100R002C70SPC100和客户端的安装, 包含Spark2x组件。
- 参考<http://zeppelin.apache.org/docs/latest/interpreter/spark.html>

操作步骤

- 编辑zeppelin-env.sh文件, 位置 `/opt/zeppelin-0.7.3/conf`, 加入以下内容

```
export MASTER=yarn-client
export SPARK_HOME=/opt/hadoopclient/Spark2x/spark
export HADOOP_CONF_DIR=/opt/hadoopclient/HDFS/hadoop/etc/hadoop
```

- 登陆Zeppelin, 选择右上角菜单中的 Interpreter



- 选择Spark, 点击 **edit** 编辑, 将 Master 参数改为 yarn-client, 点击 **save** 保存

Properties

name	value
args	
master	yarn-client
spark.app.name	Zeppelin

- 重启zeppelin

```
source /opt/hadoopclient/bigdata_env
kinit -kt /opt/user.keytab test
cd /opt/zeppelin-0.7.3/bin
./zeppelin-daemon.sh restart
```

- 执行zeppelin的sparkSQL语句

%spark.sql

FINISHED

▶

⌵

📖

⚙️

```
select * from employees_hive where id =2
```

📊

📈

📉

📉

📈

📉

📉

⌵

id	name	usd_flag	salary	address	entrytime
2	Tom	D	12000.02	America:NewYork	2014

- 执行zeppelin的spark样例代码zeppelin Tutorial -> Basic Features(Spark)

样例代码需要访问Internet上的资源，所以保证zeppelin所在的节点可以联网，检测是否能打开以下链接

Load data into table

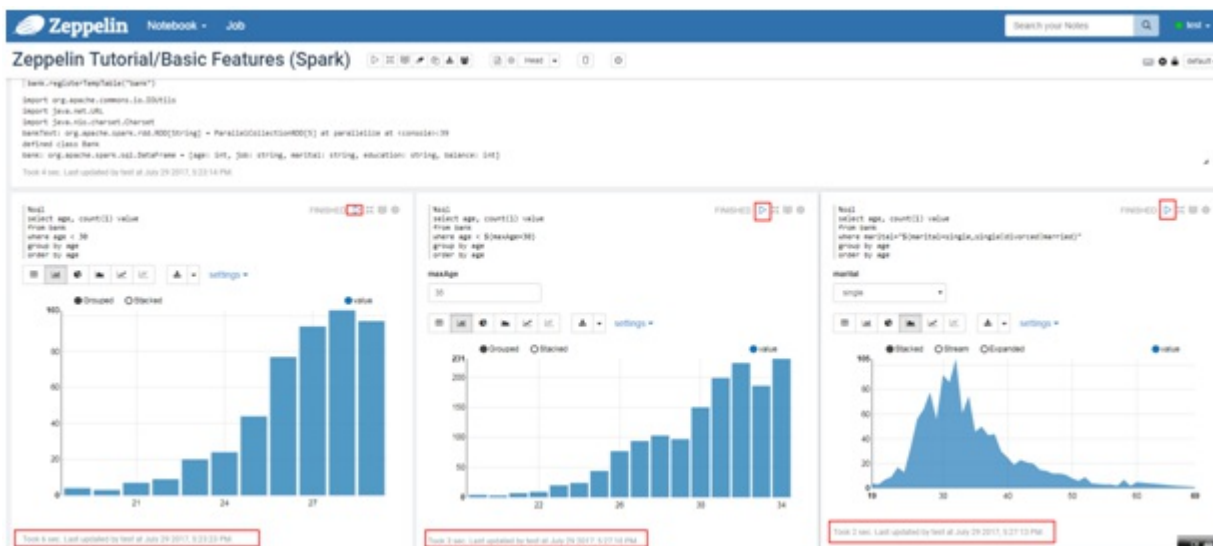
```
import org.apache.commons.io.IOUtils
import java.net.URL
import java.nio.charset.Charset

// Zeppelin creates and injects sc (SparkContext) and sqlContext (HiveContext or SqlContext)
// So you don't need create them manually

// load bank data
val bankText = sc.parallelize(
  IOUtils.toString(
    new URL("https://s3.amazonaws.com/apache-zeppelin/tutorial/bank/bank.csv"),
    Charset.forName("utf8")).split("\n"))

case class Bank(age: Integer, job: String, marital: String, education: String, balance: Integer)

val bank = bankText.map(s => s.split(";")).filter(s => s(0) != "\"age\"").map(
  s => Bank(s(0).toInt,
    s(1).replaceAll("\"", ""),
    s(2).replaceAll("\"", ""),
    s(3).replaceAll("\"", ""),
    s(5).replaceAll("\"", ").toInt
)
).toDF()
bank.registerTempTable("bank")
```



- 执行zeppelin的spark样例代码Zeppelin Tutorial -> Matplotlib (Python • PySpark)

安装python-matplotlib

```
yum install python-matplotlib
```

安装Anaconda2-4.4

```
wget https://repo.continuum.io/archive/Anaconda2-4.4.0-Linux-x86_64.sh
sh Anaconda2-4.4.0-Linux-x86_64.sh
```

配置环境变量PATH，将python换成安装Anaconda安装目录中的python

```
export PATH=/root/anaconda2/bin/:$PATH
```

在zeppelin的界面中，选择右上角的 Interpreter

选择Spark，点击 **edit** 编辑，将 zeppelin.pyspark.python 参数改为Anaconda安装目录中的python，点击 **save** 保存

zeppelin.pyspark.python


/root/anaconda2/bin/python

执行zeppelin的pyspark样例代码Zeppelin Tutorial -> Matplotlib

First line (figure will be displayed here)

```
%pyspark
import matplotlib.pyplot as plt
plt.close() # Added here to reset the plot when rerunning the paragraph
%configure_mpl(angular=True, close=False)
plt.plot([1, 2, 3], label=r'$y=x$')
[<matplotlib.lines.Line2D object at 0x7fa9c369ef00>]
```

FINISHED ▶ ⚙ ⌂



Took 30 sec. Last updated by test at July 31 2017, 5:15:10 PM

Zeppelin连接SparkR

操作场景

Zeppelin中配置Spark解析器，连接SparkR

前提条件

- 完成Zeppelin0.7.3的安装;
- 已完成FusionInsight HD V100R002C70SPC100和客户端的安装，包含Spark组件。
- 参考<http://zeppelin.apache.org/docs/latest/interpreter/spark.html>

操作步骤

- 由于Spark的Executor上也需要执行R，所以除了在Zeppelin的节点上安装R以外，所有FusionInsight集群节点上也要安装同版本的R，安装步骤如下：

不同OS配置yum源时下载的文件路径有所不同，下面以Redhat6.6安装R为例

如果安装R的节点无法访问互联网，参考FAQ进行R的安装

- 配置Redhat6.6的yum源

```
cd ~
rpm -aq | grep yum | xargs rpm -e --nodeps
wget http://mirrors.163.com/centos/6/os/x86_64/Packages/python-iniparse-0.3.1-2.1.el6.noarch.rpm
wget http://mirrors.163.com/centos/6/os/x86_64/Packages/yum-metadata-parser-1.1.2-16.el6.x86_64.rpm
wget http://mirrors.163.com/centos/6/os/x86_64/Packages/yum-3.2.29-81.el6.centos.noarch.rpm
wget http://mirrors.163.com/centos/6/os/x86_64/Packages/yum-plugin-fastestmirror-1.1.30-40.el6.noarch.rpm
wget http://mirrors.163.com/centos/6/os/x86_64/Packages/python-urlgrabber-3.9.1-11.el6.noarch.rpm
rpm -ivh python-iniparse-0.3.1-2.1.el6.noarch.rpm
rpm -ivh yum-metadata-parser-1.1.2-16.el6.x86_64.rpm
rpm -U python-urlgrabber-3.9.1-11.el6.noarch.rpm
rpm -ivh yum-3.2.29-81.el6.centos.noarch.rpm yum-plugin-fastestmirror-1.1.30-40.el6.noarch.rpm
cd /etc/yum.repos.d/
wget http://mirrors.163.com/.help/CentOS6-Base-163.repo
sed -i 's/$releasever/6/g' /etc/yum.repos.d/CentOS6-Base-163.repo
yum clean all
yum makecache
```

- 配置EPEL的源

Redhat 6.x 使用下面命令安装

```
rpm -Uvh https://mirrors.tuna.tsinghua.edu.cn/epel//6/x86_64/epel-release-6-8.noarch.rpm
```

- 更新cache

```
yum clean all  
yum makecache
```

- 执行 `yum install R` 安装R的相关的包
- 执行 `R`，检查R是否可用


正常启动如下图所示

```
[root@work opt]# R  
  
R version 3.4.1 (2017-06-30) -- "Single Candle"  
Copyright (C) 2017 The R Foundation for Statistical Computing  
Platform: x86_64-pc-linux-gnu (64-bit)  
  
R is free software and comes with ABSOLUTELY NO WARRANTY.  
You are welcome to redistribute it under certain conditions.  
Type 'license()' or 'licence()' for distribution details.  
  
Natural language support but running in an English locale  
  
R is a collaborative project with many contributors.  
Type 'contributors()' for more information and  
'citation()' on how to cite R or R packages in publications.  
  
Type 'demo()' for some demos, 'help()' for on-line help, or  
'help.start()' for an HTML browser interface to help.  
Type 'q()' to quit R.  
  
>
```

- FusionInsight客户端下测试是否可以使用sparkR

```
source /opt/hadoopclient/bigdata_env  
kinit test  
sparkR
```

- 正常启动如下图所示

 **Zeppelin** Notebook - Job

Zeppelin Tutorial/R (SparkR)

▶

⌂

📄

✍

📁

📥

👤

🔍

⌂

Head

▼

🗑

⌂

Hello R

```
%r
foo <- TRUE
print(foo)
bare <- c(1, 2.5, 4)
print(bare)
double <- 15.0
print(double)
```

[1] TRUE

[1] 1.0 2.5 4.0

[1] 15

Took 0 sec. Last updated by test at August 07 2017, 4:32:26 PM.

FAQ

- FusionInsight集群不允许访问网络，如何安装R**

在集群外同版本的Redhat版本下按照本文中yum源的方式进行安装R的操作，最后一步不要执行 `yum install R`

执行 `yum install yum-utils` 安装yumdownloader

执行 `yumdownloader R --resolve --destdir=/tmp/packages` 把所有的rpm安装包下载到 `/tmp/packages` 中

将 `/tmp/packages` 中的所有rpm包复制到集群每个节点的 `/tmp/packages` 中

切换到集群每个节点的 `/tmp/packages` 中，执行 `yum localinstall *.rpm` 完成安装
- 连接hbase出现AuthFialed for /hwbakup/hbase**

```
%hbase
help

org.apache.zookeeper KeeperException$AuthFailedException: KeeperErrorCode = AuthFailed for /hwbakup/hbase
    at org.apache.zookeeper KeeperException.create(KeeperException.java:123)
    at org.apache.zookeeper KeeperException.create(KeeperException.java:51)
    at org.apache.zookeeper ZooKeeper.create(ZooKeeper.java:783)
    at org.apache.hadoop.hbase.zookeeper.RecoverableZooKeeper.createNonSequential(RecoverableZooKeeper.java:512)
    at org.apache.hadoop.hbase.zookeeper.RecoverableZooKeeper.create(RecoverableZooKeeper.java:491)
    at org.apache.hadoop.hbase.zookeeper.ZKUtil.createWithParents(ZKUtil.java:1252)
    at org.apache.hadoop.hbase.zookeeper.ZKUtil.createWithParents(ZKUtil.java:1230)
    at com.huawei.hadoop.hbase.backup.zookeeper.BackupZooKeeperWatcher.createBaseZNodes(BackupZooKeeperWatcher.java:137)
    at com.huawei.hadoop.hbase.backup.zookeeper.BackupZooKeeperWatcher.<init>(BackupZooKeeperWatcher.java:78)
    at com.huawei.hadoop.hbase.backup.client.BackupAdmin.<init>(BackupAdmin.java:98)
    at sun.reflect.NativeConstructorAccessorImpl.newInstance(Native Method)
    at sun.reflect.NativeConstructorAccessorImpl.newInstance(NativeConstructorAccessorImpl.java:57)
    at sun.reflect.DelegatingConstructorAccessorImpl.newInstance(DelegatingConstructorAccessorImpl.java:45)
    at java.lang.reflect.Constructor.newInstance(Constructor.java:526)
    at org.jruby.javasupport.JavaConstructor.newInstanceDirect(JavaConstructor.java:275)
```

原因: zeppelin的原理hbase的jar包与从FusionInsight客户端下拷贝过来的jar冲突。

解决: 将zeppelin中原有的重名jar包移走或删除，全部用FusionInsight客户端下的相关jar包。

- Zeppelin连接spark是报如下 NoSuchMethodError**

```
java.lang.NoSuchMethodError: scala.reflect.api.JavaUniverse.runtimeMirror(Ljava/lang/ClassLoader;)Lscala/reflect/api/JavaMirrors$JavaMirror;
    at org.apache.spark.repl.SparkILoop.<init>(SparkILoop.scala:936)
    at org.apache.spark.repl.SparkILoop.<init>(SparkILoop.scala:70)
    at org.apache.zeppelin.spark.SparkInterpreter.open(SparkInterpreter.java:775)
    at org.apache.zeppelin.interpreter.LazyOpenInterpreter.open(LazyOpenInterpreter.java:70)
    at org.apache.zeppelin.spark.PySparkInterpreter.getSparkInterpreter(PySparkInterpreter.java:564)
    at org.apache.zeppelin.spark.PySparkInterpreter.createGatewayServerAndStartScript(PySparkInterpreter.java:208)
    at org.apache.zeppelin.spark.PySparkInterpreter.open(PySparkInterpreter.java:162)
    at org.apache.zeppelin.interpreter.LazyOpenInterpreter.open(LazyOpenInterpreter.java:70)
    at org.apache.zeppelin.interpreter.remote.RemoteInterpreterServer$InterpreterJob.jobRun(RemoteInterpreterServer.java:491)
    at org.apache.zeppelin.scheduler.Job.run(Job.java:175)
    at org.apache.zeppelin.scheduler.FIFOscheduler$1.run(FIFOscheduler.java:139)
    at java.util.concurrent.Executors$RunnableAdapter.call(Executors.java:471)
    at java.util.concurrent.FutureTask.run(FutureTask.java:262)
    at java.util.concurrent.ScheduledThreadPoolExecutor$ScheduledFutureTask.access$201(ScheduledThreadPoolExecutor.java:178)
    at java.util.concurrent.ScheduledThreadPoolExecutor$ScheduledFutureTask.run(ScheduledThreadPoolExecutor.java:292)
    at java.util.concurrent.ThreadPoolExecutor.runWorker(ThreadPoolExecutor.java:1145)
    at java.util.concurrent.ThreadPoolExecutor$Worker.run(ThreadPoolExecutor.java:615)
```

原因: jar包冲突

解决: 删除 `/opt/zeppelin-0.7.3/lib/` 下原有jar包scala-reflect-2.11.7.jar，替换为FusionInsight客户端下的jar包，重启zeppelin

- Zeppelin执行Spark样例代码时报GC overhead limit exceeded

```
Load data into table

import org.apache.commons.io.IOUtils
import java.net.URL
import java.nio.charset.Charset

// Zeppelin creates and injects sc (SparkContext) and sqlContext (HiveContext or SqlContext)
// So you don't need create them manually

// load bank data
val bankText = sc.parallelize(
  IOUtils.toString(
    new URL("https://s3.amazonaws.com/apache-zeppelin/tutorial/bank/bank.csv"),
    Charset.forName("utf8")).split("\n"))

case class Bank(age: Integer, job: String, marital: String, education: String, balance: Integer)

val bank = bankText.map(s => s.split(";")).filter(s => s(0) != "\age\").map(
  s => Bank(s(0).toInt,
    s(1).replaceAll("\", ""),
    s(2).replaceAll("\", ""),
    s(3).replaceAll("\", ""),
    s(5).replaceAll("\", ").toInt
  )
).toDF()
bank.registerTempTable("bank")

java.lang.OutOfMemoryError: GC overhead limit exceeded
```

原因：内存不够

解决：安装Zeppelin的节点的内存需要16G以上

- 执行zeppelin的样例代码Zeppelin Tutorial/Matplotlib (Python PySpark)报如下错误

```
Add title ERROR ▶ ✕ ⓘ ⚙

%pyspark
plt.title('Inline plotting example', fontsize=20)

<matplotlib.text.Text object at 0x2af1d90>

Traceback (most recent call last):
  File "/tmp/zeppelin_pyspark-6355775624574224283.py", line 367, in <module>
    raise Exception(traceback.format_exc())
Exception: Traceback (most recent call last):
  File "/tmp/zeppelin_pyspark-6355775624574224283.py", line 365, in <module>
    exec(code, _zcUserQueryNamespace)
  File "<stdin>", line 2, in <module>
  File "/opt/zeppelin-0.7.2-bin-all/interpreter/lib/python/backend_zinline.py", line 303, in displayhook
    show()
  File "/opt/zeppelin-0.7.2-bin-all/interpreter/lib/python/backend_zinline.py", line 72, in __call__
    manager.show(**kwargs)
  File "/opt/zeppelin-0.7.2-bin-all/interpreter/lib/python/backend_zinline.py", line 210, in show
    self.canvas.draw_idle()
  File "/opt/zeppelin-0.7.2-bin-all/interpreter/lib/python/backend_zinline.py", line 134, in draw_idle
    if not self._is_idle_drawing:
AttributeError: 'FigureCanvasZInline' object has no attribute '_is_idle_drawing'

Took 0 sec. Last updated by test at August 01 2017, 8:17:53 PM.
```

原因：python版本问题

解决：安装Anaconda2-4.4